



# Parallel Computing with MATLAB



**Sarah Wait Zaraneek, PhD**

**Senior Application Engineer**

**Jamie Winter**

**Senior Account Manager**

## Some Questions to Consider

- Do you want to speed up your algorithms?
- Do you have datasets too big to fit on your computer?

If so...

- Do you have a multicore or multiprocessor desktop machine?
- Do you have access to a computer cluster?

# Solving Big Technical Problems

## Challenges

Long running

---

Computationally  
intensive

## You could...

Wait



## Solutions

Run similar *tasks*  
on independent  
processors in  
*parallel*



Large data set

Reduce size  
of problem

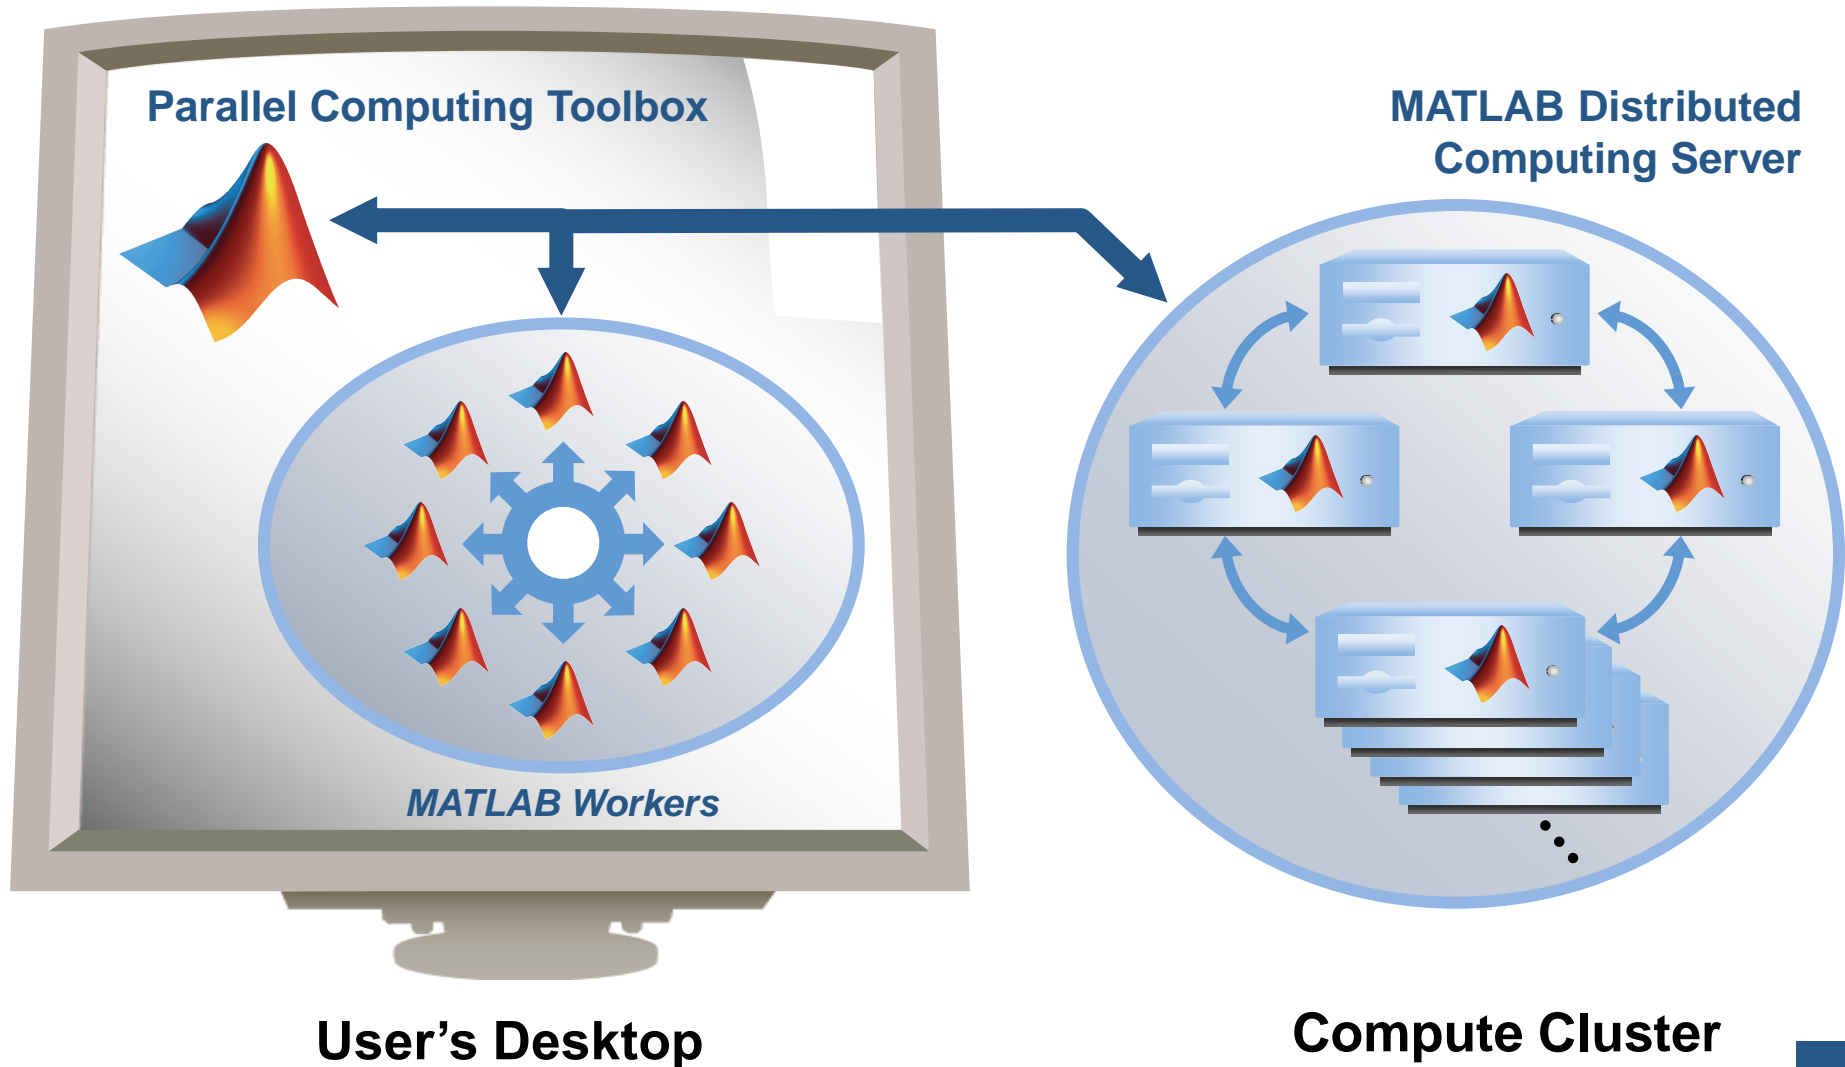


Load *data* onto  
multiple machines  
that work together  
in *parallel*

# Parallel Computing with MATLAB



# Parallel Computing with MATLAB



# Programming Parallel Applications

*Level of control*

**Minimal**

**Some**

**Extensive**

*Level of effort*

**None**

**Straightforward**

**Involved**



# Parallel Computing with MATLAB

*Level of effort*

**None**

- Built-in Toolbox Support

**Straightforward**

*Task Parallel*



*Data Parallel*

- `parfor`
- `job` and `tasks`
- `distributed`
- `spmd`
- MATLAB and MPI

**Involved**

# Example: Optimizing Tower Placement



- Determine location of cell towers
- Maximize coverage
- Minimize overlap





# Summary of Example

- Enabled built-in support for Parallel Computing Toolbox in Optimization Toolbox
- Used a pool of MATLAB workers
- Optimized in parallel using `fmincon`




# Parallel Computing Support in Optimization Toolbox

- Functions:
  - **fmincon**
    - Finds a constrained minimum of a function of several variables
  - **fminimax**
    - Finds a minimax solution of a function of several variables
  - **fgoalattain**
    - Solves the multiobjective goal attainment optimization problem
- Functions can take finite differences in parallel in order to speed the estimation of gradients



# Tools Providing Parallel Computing Support

- Optimization Toolbox
- GADS Toolbox
- Statistics Toolbox
- SystemTest
- Simulink Design Optimization
- Bioinformatics Toolbox
- Model-Based Calibration Toolbox
- ...



*Directly leverage functions in Parallel Computing Toolbox*

# Task Parallel Applications





# Example: Parameter Sweep of ODEs

- Solve a 2<sup>nd</sup> order ODE



$$\overbrace{m}^5 \ddot{x} + \underbrace{b}_{1,2,\dots} \dot{x} + \underbrace{k}_{1,2,\dots} x = 0$$

- Simulate with different values for **b** and **k**
- Record peak value for each run
- Plot results

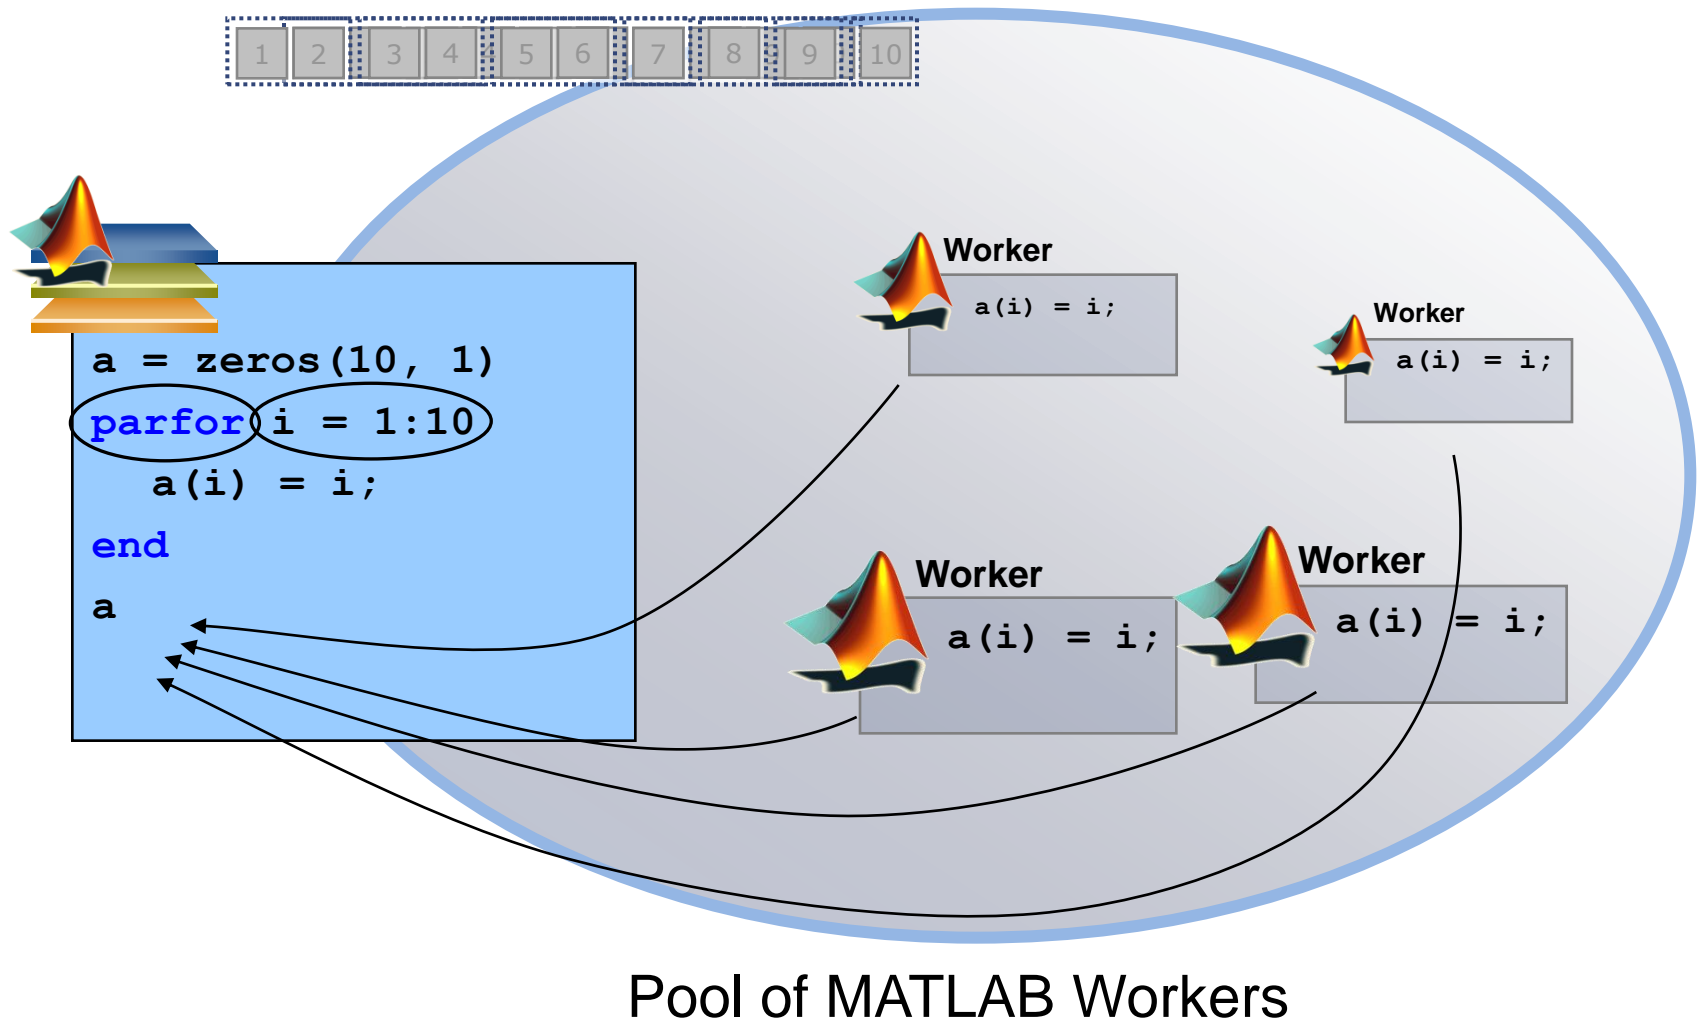


# Summary of Example

- Mixed task-parallel and serial code in the same function
- Ran loops on a pool of MATLAB resources
- Used M-Lint analysis to help in converting existing `for`-loop into `parfor`-loop



# The Mechanics of parfor Loops



# Converting `for` to `parfor`

- Requirements for `parfor` loops
  - Task independent
  - Order independent
- Constraints on the loop body
  - Cannot “introduce” variables (e.g. `eval`, `load`, `global`, etc.)
  - Cannot contain `break` or `return` statements
  - Cannot contain another `parfor` loop



# Advice for Converting `for` to `parfor`

- Use M-Lint to diagnose `parfor` issues
- If your `for` loop cannot be converted to a `parfor`, consider wrapping a subset of the body to a function
- Read the section in the documentation on classification of variables
- <http://blogs.mathworks.com/loren/2009/10/02/using-parfor-loops-getting-up-and-running/>

# Parallel Computing Tools Address...

## Task-Parallel

### Long computations

- Multiple independent iterations

```
parfor i = 1 : n
    % do something with i
end
```

- Series of tasks

Task 1

Task 2

Task 3

Task 4


## Data-Parallel

### Large data problems

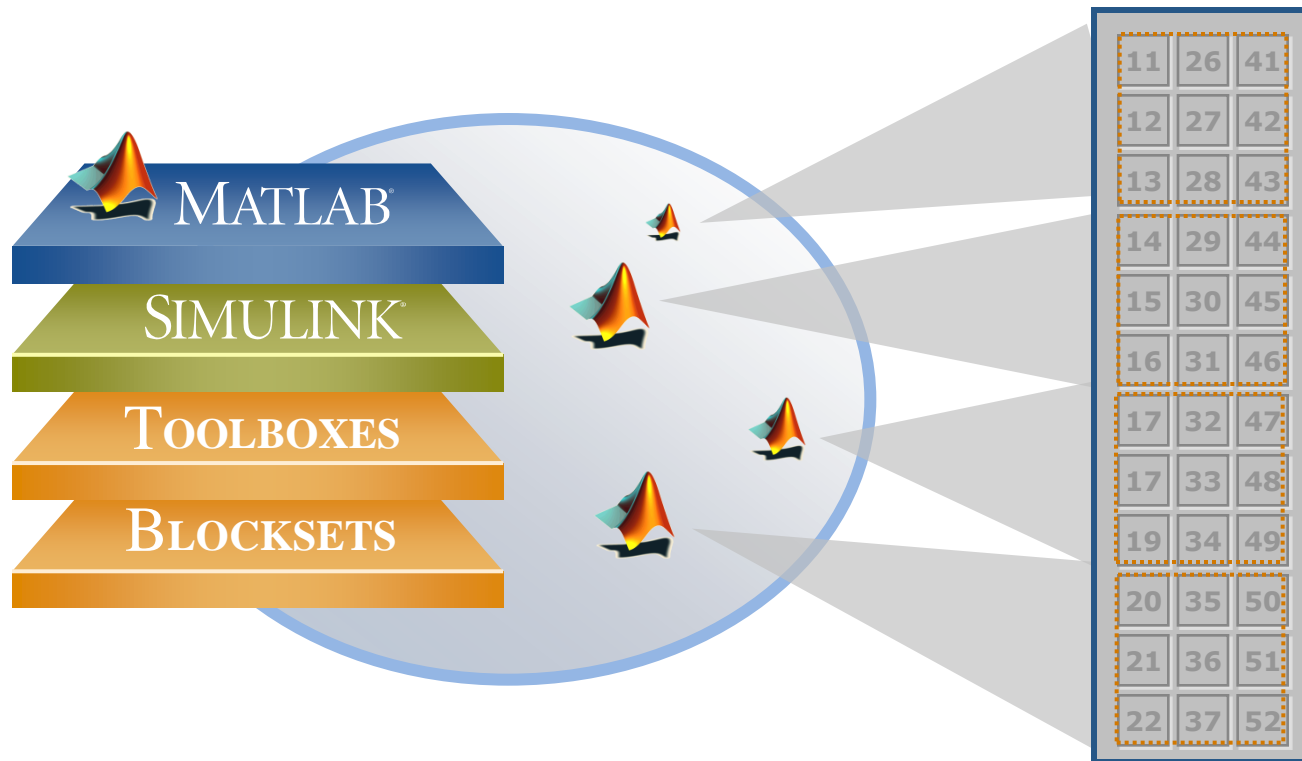
11	26	41
12	27	42
13	28	43
14	29	44
15	30	45
16	31	46
17	32	47
17	33	48
19	34	49
20	35	50
21	36	51
22	37	52



# Large Data Sets (Data Parallel)



# Client-side Distributed Arrays



Remotely Manipulate Array  
from Desktop

Distributed Array  
Lives on the Cluster

# Distributed Arrays and Parallel Algorithms

- Distributed arrays
  - Store segments of data across participating workers
  - Can be created and manipulated directly from the client.
  - Create from any built-in class in MATLAB
    - Examples: doubles, sparse, logicals, cell arrays, and arrays of structs
- Parallel algorithms for distributed arrays
  - Matrix manipulation operations
    - Examples: indexing, data type conversion, and transpose
  - Parallel linear algebra functions, such as **svd** and **lu**

# Enhanced MATLAB Functions That Operate on Distributed Arrays

Type of Function	Function Names
Data functions	<a href="#">cumprod</a> , <a href="#">cumsum</a> , <a href="#">fft</a> , <a href="#">max</a> , <a href="#">min</a> , <a href="#">prod</a> , <a href="#">sum</a>
Data type functions	<a href="#">cast</a> , <a href="#">cell2mat</a> , <a href="#">cell2struct</a> , <a href="#">celldisp</a> , <a href="#">cellfun</a> , <a href="#">char</a> , <a href="#">double</a> , <a href="#">fieldnames</a> , <a href="#">int16</a> , <a href="#">int32</a> , <a href="#">int64</a> , <a href="#">int8</a> , <a href="#">logical</a> , <a href="#">num2cell</a> , <a href="#">rmfield</a> , <a href="#">single</a> , <a href="#">struct2cell</a> , <a href="#">swapbytes</a> , <a href="#">typecast</a> , <a href="#">uint16</a> , <a href="#">uint32</a> , <a href="#">uint64</a> , <a href="#">uint8</a>
Elementary and trigonometric functions	<a href="#">abs</a> , <a href="#">acos</a> , <a href="#">acosh</a> , <a href="#">acot</a> , <a href="#">acotd</a> , <a href="#">acoth</a> , <a href="#">acsc</a> , <a href="#">acscd</a> , <a href="#">acsch</a> , <a href="#">angle</a> , <a href="#">asec</a> , <a href="#">asecd</a> , <a href="#">asech</a> , <a href="#">asin</a> , <a href="#">asind</a> , <a href="#">asinh</a> , <a href="#">atan</a> , <a href="#">atan2</a> , <a href="#">atand</a> , <a href="#">atanh</a> , <a href="#">ceil</a> , <a href="#">complex</a> , <a href="#">conj</a> , <a href="#">cos</a> , <a href="#">cosd</a> , <a href="#">cosh</a> , <a href="#">cot</a> , <a href="#">cotd</a> , <a href="#">coth</a> , <a href="#">csc</a> , <a href="#">cscd</a> , <a href="#">csch</a> , <a href="#">exp</a> , <a href="#">expm1</a> , <a href="#">fix</a> , <a href="#">floor</a> , <a href="#">hypot</a> , <a href="#">imag</a> , <a href="#">isreal</a> , <a href="#">log</a> , <a href="#">log10</a> , <a href="#">loglp</a> , <a href="#">log2</a> , <a href="#">mod</a> , <a href="#">nextpow2</a> , <a href="#">nthroot</a> , <a href="#">pow2</a> , <a href="#">real</a> , <a href="#">reallog</a> , <a href="#">realpow</a> , <a href="#">realsqrt</a> , <a href="#">rem</a> , <a href="#">round</a> , <a href="#">sec</a> , <a href="#">secd</a> , <a href="#">sech</a> , <a href="#">sign</a> , <a href="#">sin</a> , <a href="#">sind</a> , <a href="#">sinh</a> , <a href="#">sqrt</a> , <a href="#">tan</a> , <a href="#">tand</a> , <a href="#">tanh</a>
Elementary matrices	<a href="#">cat</a> , <a href="#">diag</a> , <a href="#">eps</a> , <a href="#">find</a> , <a href="#">isempty</a> , <a href="#">isequal</a> , <a href="#">isequalwithhequalnans</a> , <a href="#">isfinite</a> , <a href="#">isinf</a> , <a href="#">isnan</a> , <a href="#">length</a> , <a href="#">ndims</a> , <a href="#">size</a> , <a href="#">tril</a> , <a href="#">triu</a>
Matrix functions	<a href="#">chol</a> , <a href="#">eig</a> , <a href="#">lu</a> , <a href="#">norm</a> , <a href="#">normest</a> , <a href="#">svd</a>
Array operations	<a href="#">all</a> , <a href="#">and</a> , <a href="#">any</a> , <a href="#">bitand</a> , <a href="#">bitor</a> , <a href="#">bitxor</a> , <a href="#">ctranspose</a> , <a href="#">end</a> , <a href="#">eq</a> , <a href="#">ge</a> , <a href="#">gt</a> , <a href="#">horzcat</a> , <a href="#">ldivide</a> , <a href="#">le</a> , <a href="#">lt</a> , <a href="#">minus</a> , <a href="#">mldivide</a> , <a href="#">mrdivide</a> , <a href="#">mtimes</a> , <a href="#">ne</a> , <a href="#">not</a> , <a href="#">or</a> , <a href="#">plus</a> , <a href="#">power</a> , <a href="#">rdivide</a> , <a href="#">subsasgn</a> , <a href="#">subsindex</a> , <a href="#">subsref</a> , <a href="#">times</a> , <a href="#">transpose</a> , <a href="#">uminus</a> , <a href="#">uplus</a> , <a href="#">vertcat</a> , <a href="#">xor</a>
Sparse matrix functions	<a href="#">full</a> , <a href="#">issparse</a> , <a href="#">nnz</a> , <a href="#">nonzeros</a> , <a href="#">nzmax</a> , <a href="#">sparse</a> , <a href="#">spfun</a> , <a href="#">spones</a>
Special functions	<a href="#">dot</a>

# MPI-Based Functions in Parallel Computing Toolbox™

Use when a high degree of control over parallel algorithm is required

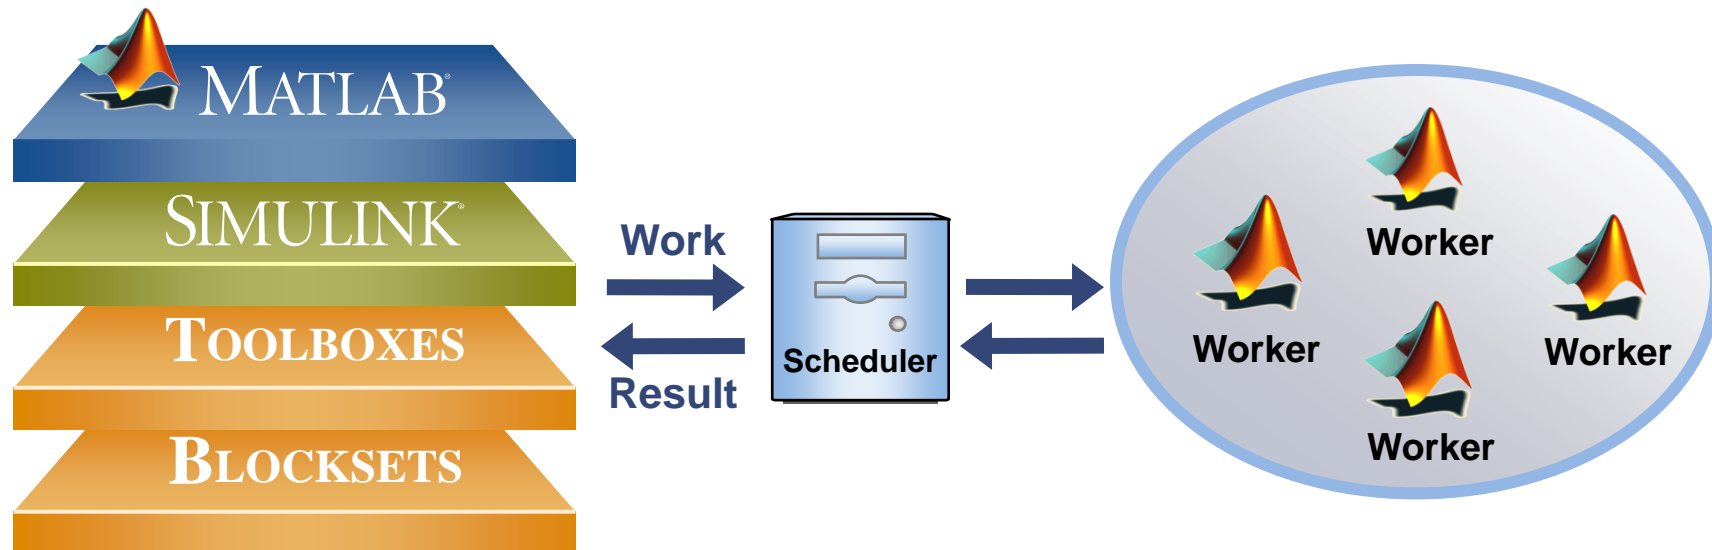
- High-level abstractions of MPI functions
  - `labSendReceive`, `labBroadcast`, and others
  - Send, receive, and broadcast any data type in MATLAB
- Automatic bookkeeping
  - Setup: communication, ranks, etc.
  - Error detection: deadlocks and miscommunications
- Pluggable
  - Use any MPI implementation that is *binary-compatible* with MPICH2

# Interactive to Scheduling

- Interactive
  - Great for prototyping
  - Immediate access to MATLAB workers
  
- Scheduling
  - Offloads work to other MATLAB workers (local or on a cluster)
  - Access to more computing resources for improved performance
  - Frees up local MATLAB session





# Scheduling Work





# Example: Schedule Processing

- Offload parameter sweep to local workers
- Get peak value results when processing is complete
- Plot results in local MATLAB



# Summary of Example

- Used `batch` for off-loading work
- Used `matlabpool` option to off-load and run in parallel
- Used `load` to retrieve worker's workspace



# Task-Parallel Workflows

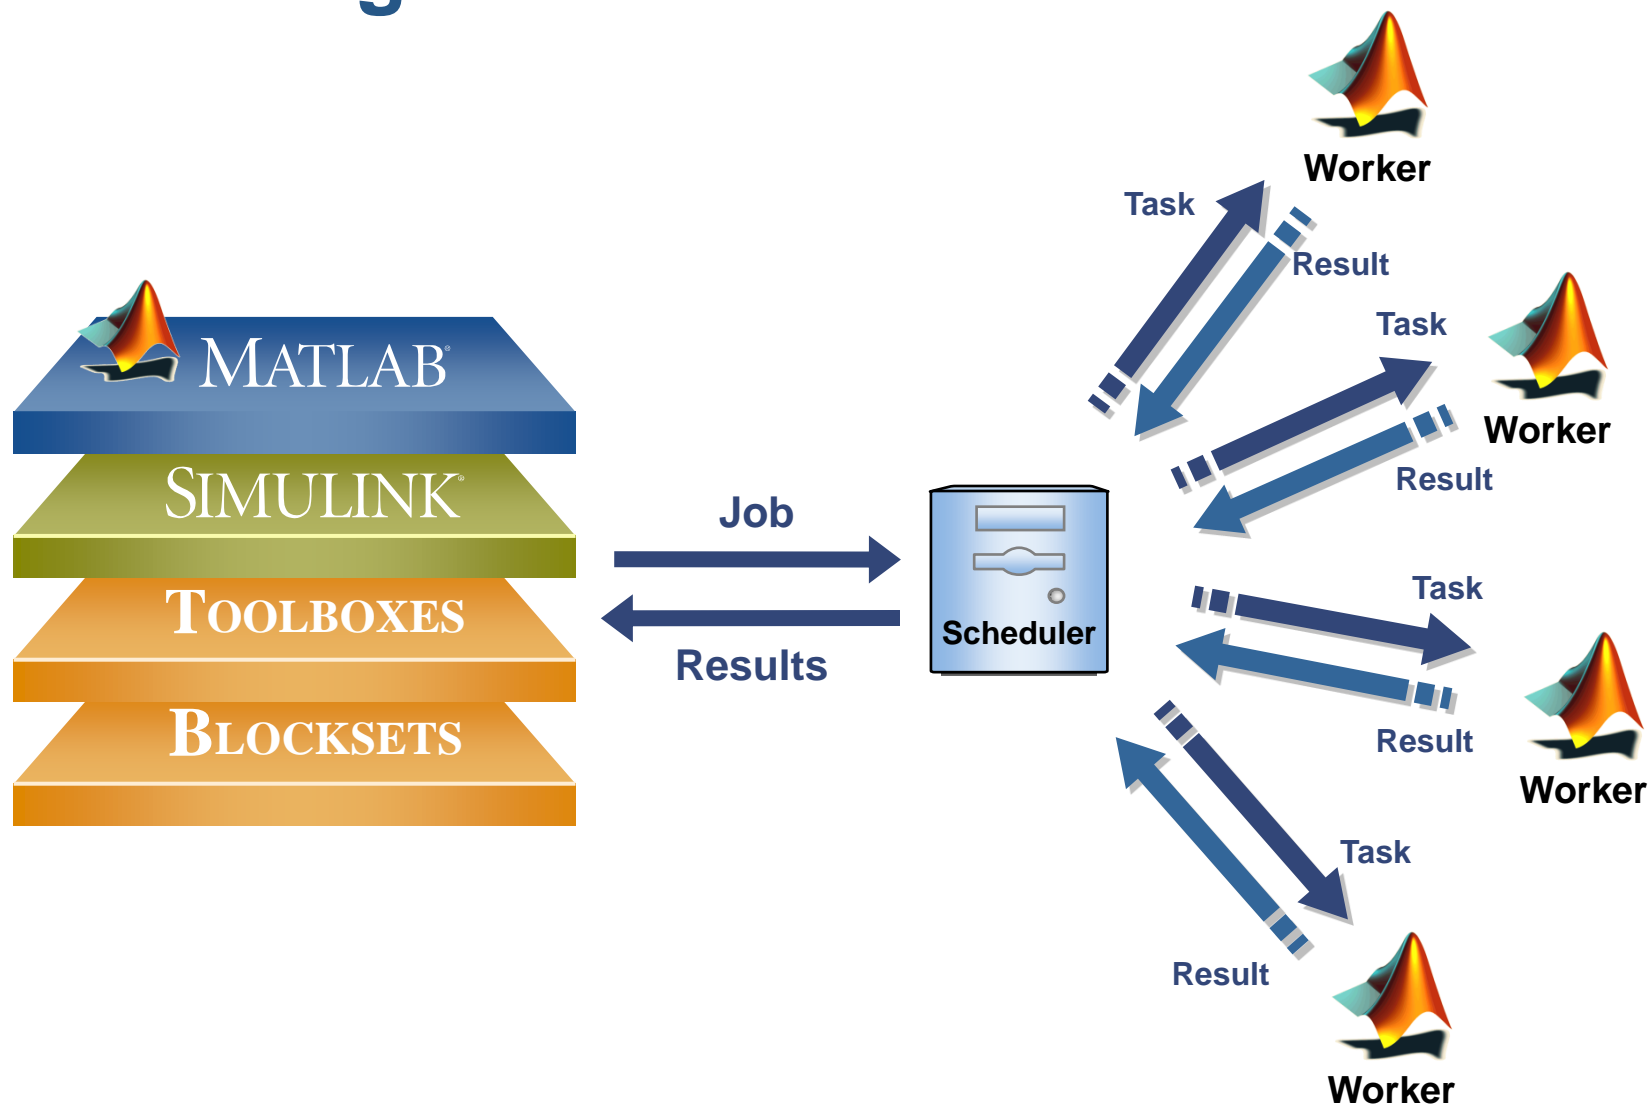
- **parfor**

- Multiple independent iterations
- Easy to combine serial and parallel code
- Workflow
  - Interactive using `matlabpool`
  - Scheduled using `batch`

- **jobs/tasks**


- Series of independent tasks; not necessarily iterations
- Workflow → Always scheduled

# Scheduling Jobs and Tasks




# Example: Scheduling Independent Simulations

- Offload three independent approaches to solving our previous ODE example
- Retrieve simulated displacement as a function of time for each simulation
- Plot comparison of results in local MATLAB



# Summary of Example

- Used `findResource` to find scheduler
- Used `createJob` and `createTask` to set up the problem
- Used `submit` to off-load and run in parallel
- Used `getAllOutputArguments` to retrieve all task outputs




# Factors to Consider for Scheduling


- There is always an overhead to distribution
  - Combine small repetitive function calls
- Share code and data with workers efficiently
  - Set job properties (`FileDependencies`, `PathDependencies`)
- Minimize I/O
  - Enable `Workspace` option for `batch`
- Capture command window output
  - Enable `CaptureDiary` option for `batch`



# Options for Scheduling Jobs



**Task Parallel**



**Data Parallel**

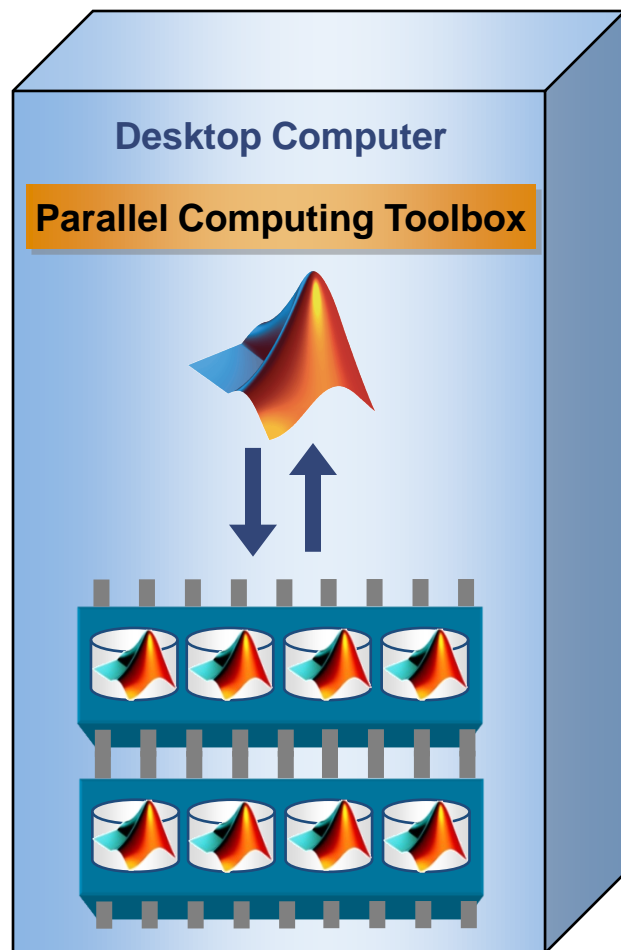
```
>> createMatlabPoolJob
or
>> batch
```

```
>> createMatlabPoolJob
or
>> batch
```

```
>> createJob (...)
>> createTask (...)
```

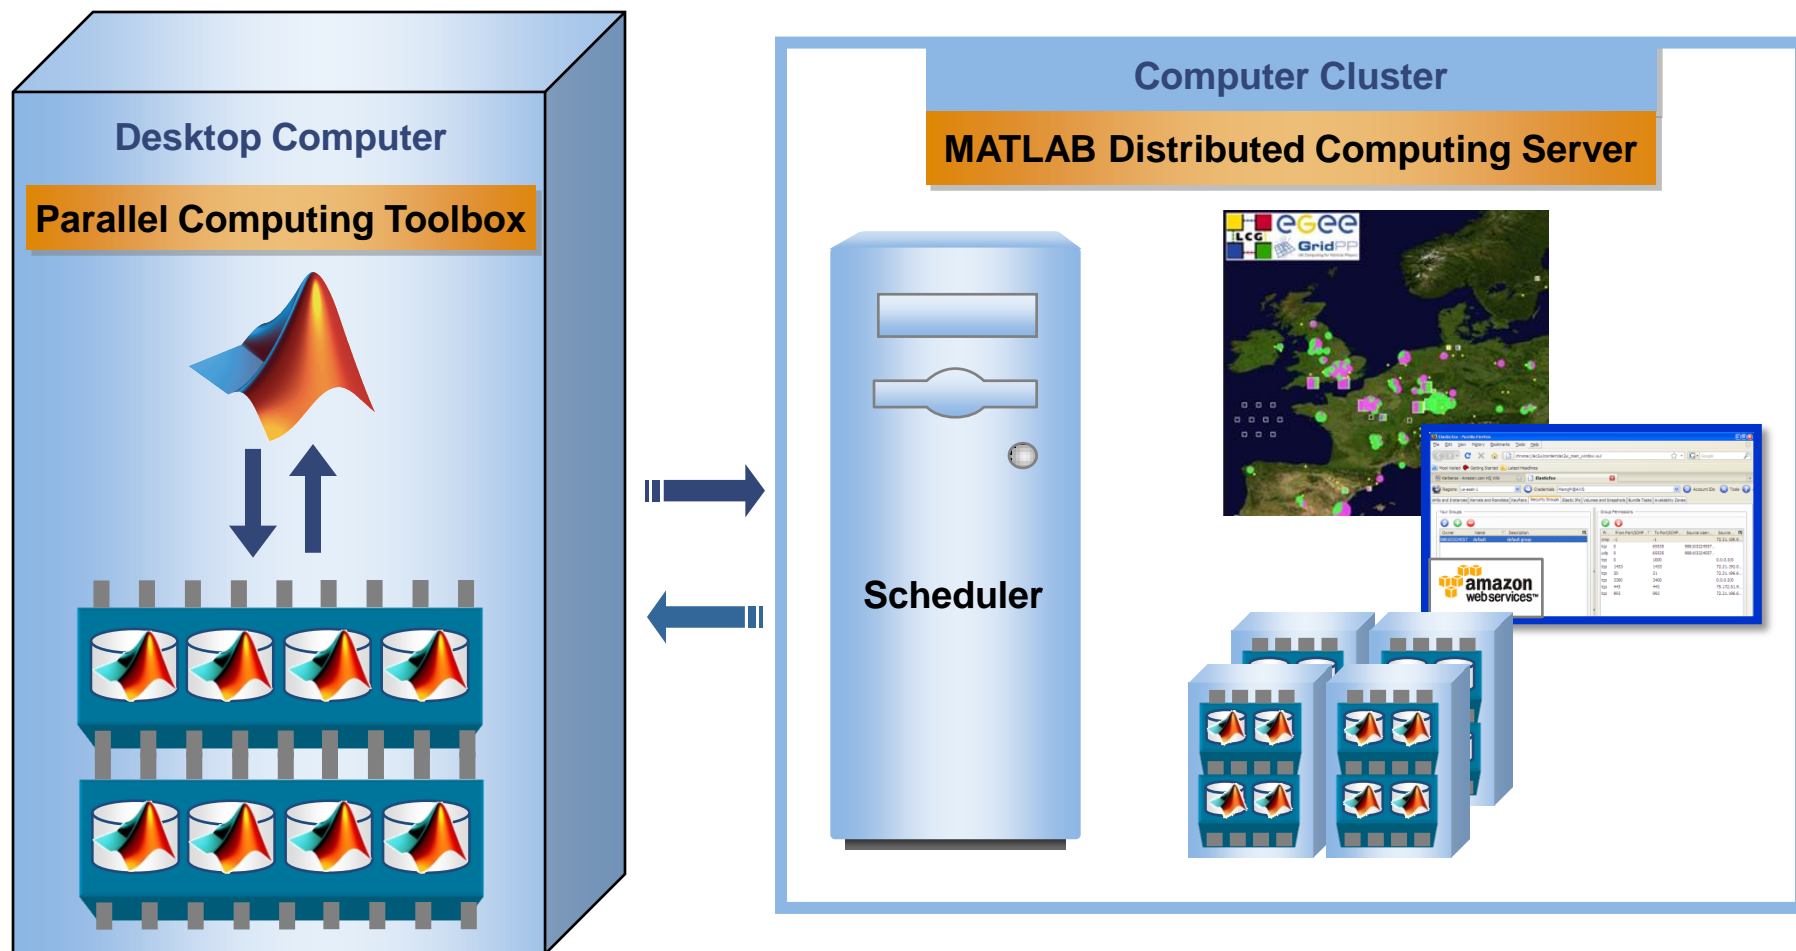
```
>> createParallelJob
```

# Run 8 Local Workers on Desktop




- Rapidly develop parallel applications on local computer
- Take full advantage of desktop power
- Separate computer cluster not required

# Scale Up to Clusters, Grids and Clouds



# Support for Schedulers

## Direct Support



**Platform™**



**TORQUE**

## Open API for others



# Summary

- Speed up algorithms without code changes
- Develop parallel code interactively
  - Task-parallel applications for faster processing
  - Data-parallel applications for handling large data sets
- Schedule your programs to run

# MathWorks Resources Available

- Stay connected – register for a MathWorks Account on [www.mathworks.com](http://www.mathworks.com)
- Join newsgroups, exchange links, and more at MATLAB Central – [www.mathworks.com/matlabcentral](http://www.mathworks.com/matlabcentral)
- Get started using MATLAB and Simulink – check out our video tutorials at [www.mathworks.com/academia/student\\_center/tutorials/launchpad.html](http://www.mathworks.com/academia/student_center/tutorials/launchpad.html)

# Do you want to know more?

More info:

Jamie.Winter@mathworks.com

(508) 647 – 7463

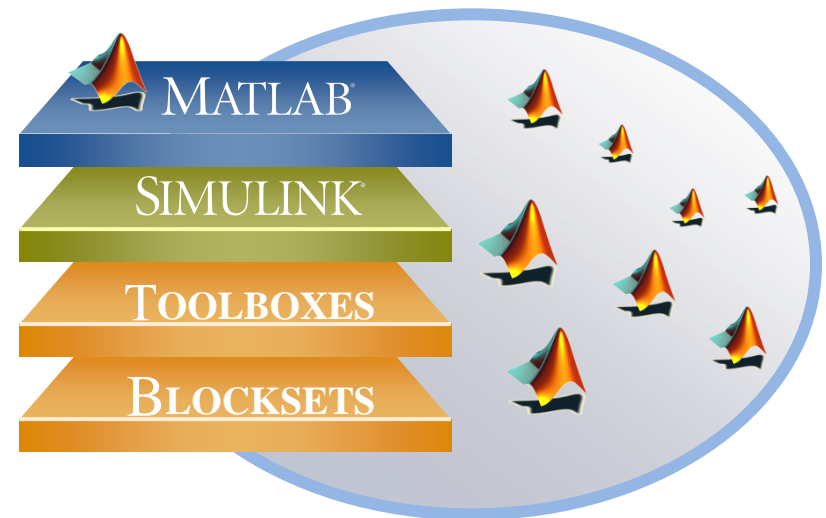






# Composite Arrays

- Created from desktop
- Stored on workers
- Syntax *similar* to cell arrays

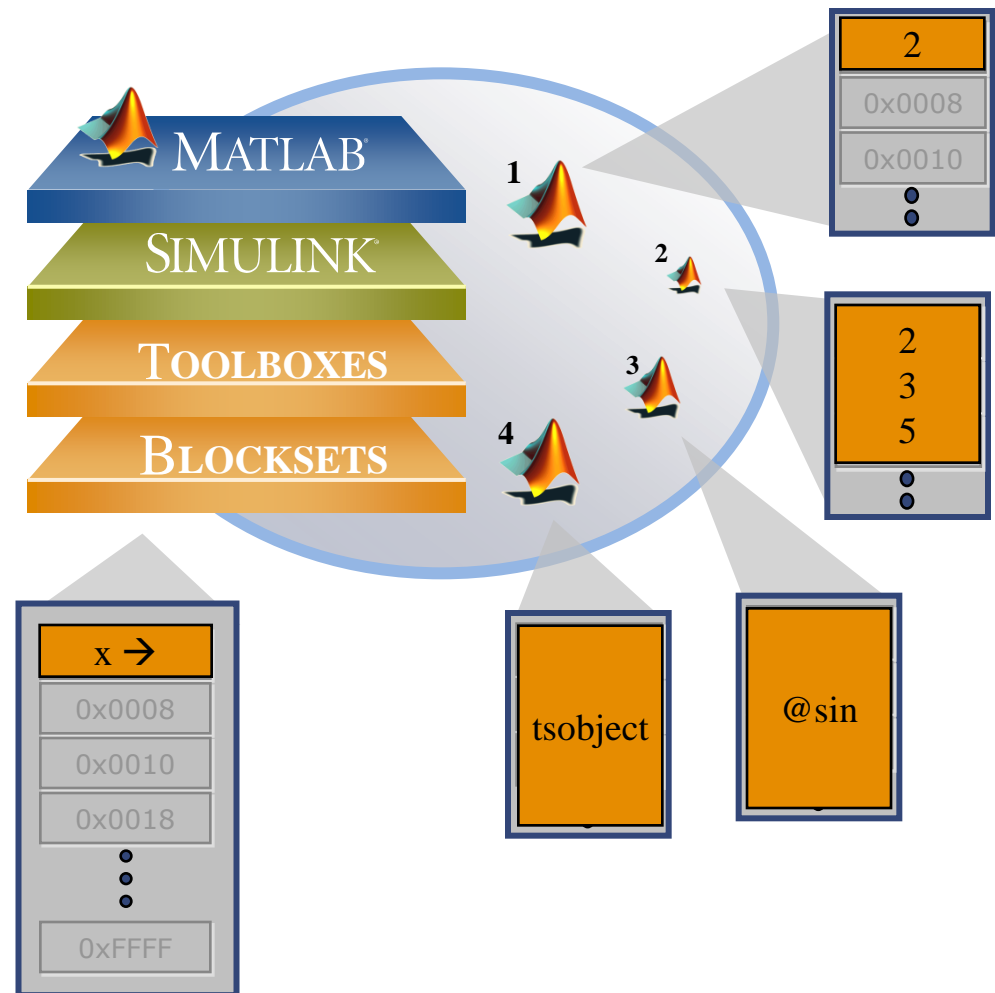


# Composite Array in Memory

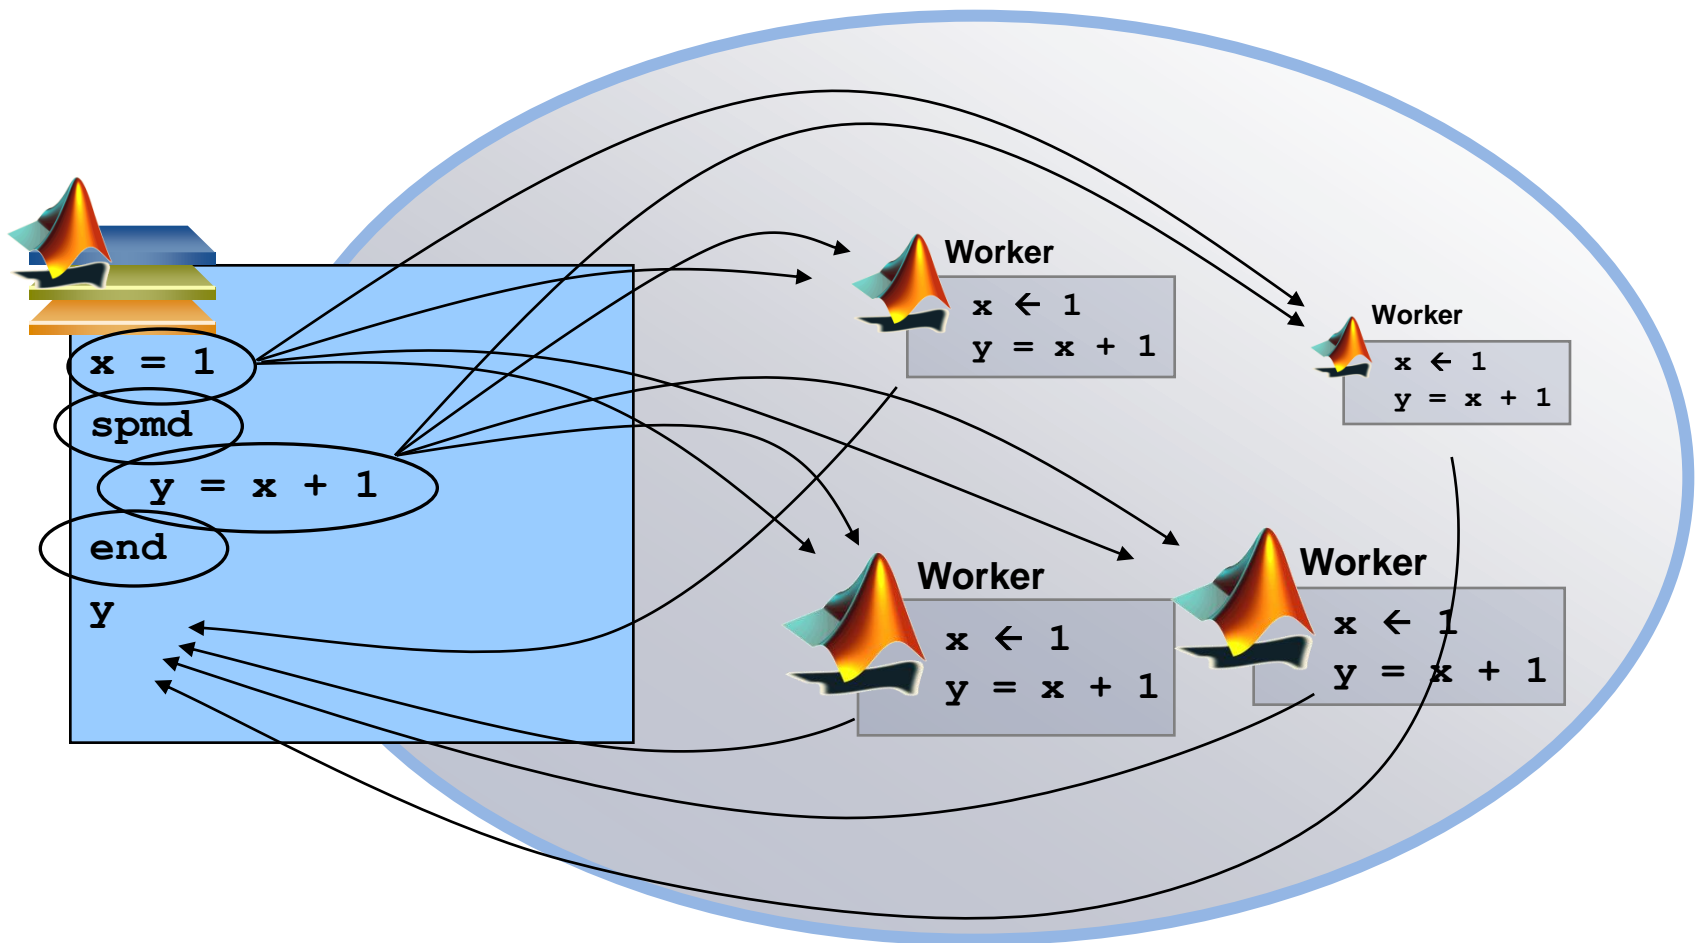
```
>> matlabpool open 4

>> x = Composite(4)

>> x{1} = 2
>> x{2} = [2, 3, 5]
>> x{3} = @sin
>> x{4} = tsobject()
```



# A mental model for SPMD ... END



Pool of MATLAB Workers